


Good to eat or deadly poison?

Machine Learning Techniques for Mushroom Classification



Introduction

- The danger of mis-identifying mushrooms: edible vs poisonous varieties.
- Motivation: use data + ML to help classify mushrooms.
- Overview: dataset → modeling → evaluation → insights & recommendations.



The Enhanced Data Set: Balanced, Categorical, and Field-Ready

Optimizing AI for Mushroom Safety through Data Structure

Dataset Balance & Size



Edible
(≈50%)


Poisonous
(≈50%)



Total Records: >25,000


Prevents model bias
towards either class

Categorical Descriptors




All 23 variables are discrete
and categorical

Field Deployment Benefits




Simplifies input; No measuring
tools required




Feature Analysis: Cramér's V Association Test


Primary Predictive Power (High Association)




Feature Interdependence (Low Cross-Correlation)



Experimental Machine Learning Pipeline: From Baselines to Ensemble Stacking



Key Performance Metrics

 Accuracy	Accuracy prioritization information specification	
 ROC-AUC	ROC-AUC metric is base recent ROC-AUC	
 F2-Score	F2-Score is inversely proportional to probability to score	
 False-Positive Rate	False-Positive rate is size and negatives	
 False-Negative Rate	False-negative rate is ensure False-negative rate	

Key Performance Metrics

Tracking comprehensive metrics to assess model reliability and risk.

Safety-First Prioritization

Safety,
Raw
Accuracy



Safety-First Prioritization

Explicitly penalizing missed toxic cases to ensure safety dominates final selection.




Table 2: All model performance summary

Model	F2	Accuracy	TPR	FNR	FPR	TNR
LDA	0.80	0.71	0.88	0.12	0.43	0.57
QDA	0.77	0.61	0.86	0.14	0.60	0.40
Random Forest	0.81	0.68	0.90	0.10	0.50	0.51
Boosting	0.82	0.70	0.89	0.10	0.46	0.54
Logit Reg.	0.82	0.71	0.89	0.12	0.44	0.56
SVM	0.76	0.73	0.79	0.21	0.32	0.68
KNN	0.79	0.60	0.89	0.11	0.64	0.36
Stack Ensemble	0.82	0.71	0.89	0.11	0.45	0.55

Stacked Model Performance Analysis & Recommendation

Evaluation of F2-score, Accuracy, and Error Rates for Mushroom Safety



F2-score
(Recall-Correctness Trade-off)

0.82

→ Best
Trade-off

Balances recall of poisonous
species and overall correctness.



Accuracy
(Overall Correctness)

0.70

Represents overall model
correctness.



False-Negative Rate
(Critical Safety Metric)

11%

Lowest Among
Contenders

Crucial for minimizing missed
poisonous species.



False-Positive Rate
(Safety "Insurance Premium")

45%

Deemed an acceptable cost for
public safety.



Key Insight: Prioritizing Public Safety

With the lowest false-negative rate and the best trade-off (F2-score), the stacked model prioritizes minimizing the risk of severe outcomes. The 45% false-positive rate is acceptable as a safety precaution.

Recommended Production Classifier



Feature Importance in Mushroom Safety Analysis: Key Predictors Ranked


Based on Performance Drop Analysis and Robust Across Random Seeds



Rankings remain consistent across random seeds, indicating robust explanatory power.

Recommendation & Future Improvements

- For better reliability: ensure dataset covers wide variety of species — avoid focusing on just few types.
- Consider combining traditional feature-based ML with image-based deep learning (when images available) for robust classification.
- Use explainability methods (e.g. feature importance, SHAP, model-agnostic analysis) to understand and trust model decisions.
- Always treat predictions as “aids”, not definitive — misclassifying a poisonous mushroom can be fatal.



Conclusion

- ML models — especially ensemble methods — can effectively classify mushrooms as edible or poisonous with high accuracy.
- Feature-based classification (using physical characteristics) is feasible and interpretable.
- Combining good data, sound ML practices, and interpretability makes such classification safer and more reliable.
- However, model predictions should not replace expert judgement — use as supportive tool only.

